

Problem 1:

$$(a). \quad y^{(1)} w^T x^{(1)} \geq 1, \quad y^{(1)} = y = -1, \quad x^{(1)} = x = \begin{pmatrix} 1 \\ 1 \end{pmatrix}$$

$$\Rightarrow y(w_1^* x_1 + w_2^* x_2) \geq 1$$

$$\Rightarrow -1(w_1^* \cdot 1 + w_2^* \cdot 1) \geq 1$$

$$\Rightarrow w_1^* + w_2^* \leq -1$$

$$\text{Also want } \min \frac{1}{2} \|w^*\|^2 = \min \left( \frac{1}{2} \sqrt{w_1^{*2} + w_2^{*2}} \right)$$

$$\text{So we have } w_1^* = w_2^* = -\frac{1}{2}$$

$$\text{Hence } \vec{w}^* = \begin{pmatrix} -1/2 \\ -1/2 \end{pmatrix}$$

$$(b). \quad x^{(1)} = \begin{pmatrix} 1 \\ 1 \end{pmatrix}, \quad x^{(2)} = \begin{pmatrix} 1 \\ 0 \end{pmatrix}, \quad y^{(1)} = 1, \quad y^{(2)} = -1$$

When  $i=1$ :

$$y^{(1)} w^{*T} x^{(1)} = 1 \cdot (w_1^* \cdot 1 + w_2^* \cdot 1) = w_1^* + w_2^* \geq 1$$

when  $i=2$ :

$$y^{(2)} w^{*T} x^{(2)} = -1 \cdot (w_1^* \cdot 1 + w_2^* \cdot 0) = -w_1^* \geq 1 \Rightarrow w_1^* \leq -1$$

$$\text{Also want to minimize } w_1^{*2} + w_2^{*2} = |w_1^*|^2 + |w_2^*|^2.$$

$$\text{Thus, we have } w_1^* = -1, \quad w_2^* = 2$$

$$\text{So } \vec{w}^* = \begin{pmatrix} -1 \\ 2 \end{pmatrix}$$

$$(c). \quad \text{Since } b \neq 0, \text{ we have } y^{(i)} \theta^T x^{(i)} = y^{(i)} (w^{*T} x^{(i)} + b^*) \geq 1$$

$$\text{And } x^{(1)} = \begin{pmatrix} 1 \\ 1 \end{pmatrix}, \quad x^{(2)} = \begin{pmatrix} 1 \\ 0 \end{pmatrix}, \quad y^{(1)} = 1, \quad y^{(2)} = -1.$$

$$\text{When } i=1: \quad 1 \cdot (w_1^* + w_2^* + b^*) = w_1^* + w_2^* + b^* \geq 1 \Rightarrow w_1^* + w_2^* \geq 1 - b^*$$

$$\text{When } i=2: \quad -1 \cdot (w_1^* + 0 + b^*) = -w_1^* - b^* \geq 1 \Rightarrow w_1^* \leq -1 - b^*$$

$$\text{When } b^* \geq 1: \quad 1 - b^* \leq 0 \text{ and } -1 - b^* \leq 0 \Rightarrow w_1^* = -1 - b^*, w_2^* = 2. \text{ So } (\vec{w}^*, b^*) = \left( \begin{pmatrix} -1-b^* \\ 2 \end{pmatrix}, b^* \right)$$

$$\text{When } 1 \geq b^* \geq -1: \quad 1 - b^* \geq 0 \text{ and } -1 - b^* \leq 0 \Rightarrow w_1^* = -1 - b^*, w_2^* = 2. \text{ So } (\vec{w}^*, b^*) = \left( \begin{pmatrix} -1-b^* \\ 2 \end{pmatrix}, b^* \right)$$

$$\text{When } -1 \geq b^*: \quad 1 - b^* \geq 0 \text{ and } -1 - b^* \geq 0 \Rightarrow w_1^* = -1 - b^*, w_2^* = 2. \text{ So } (\vec{w}^*, b^*) = \left( \begin{pmatrix} -1-b^* \\ 2 \end{pmatrix}, b^* \right)$$

$$\text{To minimize } \|w^*\|^2, \text{ we namely want to minimize } |w_1^*| = |-1 - b^*| \Rightarrow b^* = -1$$

$$\text{Hence, } (\vec{w}^*, b^*) = \left( \begin{pmatrix} 0 \\ 2 \end{pmatrix}, -1 \right)$$

## Problem 2:

$i$	Label	Hypothesis 1 (1st iteration)				Hypothesis 2 (2nd iteration)			
		$w_0$	$f_1 \equiv \text{sign}(x_1 - \underline{2})$	$f_2 \equiv \text{sign}(x_2 - \underline{6})$	$h_1 \equiv \underline{f_1}$	$w_1$	$f'_1 \equiv \text{sign}(x_1 - \underline{2})$	$f'_2 \equiv \text{sign}(x_2 - \underline{0})$	$h_2 \equiv \underline{f'_1}$
(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)
1	-	0.1	-1	-1	-1	0.0625	-1	+1	-1
2	-	0.1	-1	-1	-1	0.0625	-1	+1	-1
3	+	0.1	+1	+1	+1	0.0625	+1	+1	+1
4	+	0.1	-1	-1	-1	0.25	-1	+1	-1
5	-	0.1	-1	+1	+1	0.25	-1	+1	-1
6	-	0.1	+1	-1	-1	0.0625	+1	+1	+1
7	+	0.1	+1	+1	+1	0.0625	+1	+1	+1
8	-	0.1	-1	-1	-1	0.0625	-1	-1	-1
9	+	0.1	-1	+1	+1	0.0625	-1	+1	-1
10	+	0.1	+1	+1	+1	0.0625	+1	+1	+1

$$\uparrow$$
  
 $\xi_1 = 0.3$

$$\uparrow$$
  
 $\xi_1 = 0.2$

$$\uparrow$$
  
 $\xi_2 = 0.375$

$$\uparrow$$
  
 $\xi_2 = 0.4375$

(c).

$$\xi_1 = 0.1 + 0.1 = 0.2$$

$$\beta_1 = \frac{1}{2} \ln \left( \frac{1-0.2}{0.2} \right) = \frac{1}{2} \cdot \ln(4) = 0.6931$$

$$\text{If correct: } w_{1,i} = 0.1 \cdot \exp(-0.6931) = 0.05$$

$$\text{If incorrect: } w_{1,i} = 0.1 \cdot \exp(0.6931) = 0.2$$

Normalize;

$$\text{If correct: } w_{1,i} = \frac{0.05}{0.05 \cdot 8 + 0.2 \cdot 2} = 0.0625$$

$$\text{If incorrect: } w_{1,i} = \frac{0.2}{0.05 \cdot 8 + 0.2 \cdot 2} = 0.25$$

(d).  $\xi_2 = 0.375$

$$\beta_2 = \frac{1}{2} \ln \left( \frac{1-0.375}{0.375} \right) = \frac{1}{2} \cdot \ln\left(\frac{5}{3}\right) = 0.2554$$

Since  $\beta_2 < \beta_1$ , the final hypothesis solely depends on  $h_1$ .

$$\text{Which is: } H(x) = h_1 = (-1, -1, +1, -1, +1, -1, +1, -1, +1, +1)^T$$

## Problem 3.

1. (a).

```
#####
# functions -- evaluation
#####

def performance(y_true, y_pred, metric="accuracy"):
    """
    Calculates the performance metric based on the agreement between the
    true labels and the predicted labels.

    Parameters
    -----
    y_true -- numpy array of shape (n,), known labels
    y_pred -- numpy array of shape (n,), (continuous-valued) predictions
    metric -- string, option used to select the performance measure
              options: 'accuracy', 'f1-score', 'auROC', 'precision',
                      'sensitivity', 'specificity'

    Returns
    -----
    score -- float, performance score
    """
    # map continuous-valued predictions to binary labels
    y_label = np.sign(y_pred)
    y_label[y_label==0] = 1

    ### ===== TODO : START ===== ###
    # part 1a: compute classifier performance

    if metric == "accuracy":
        score = metrics.accuracy_score(y_true, y_pred)
    elif metric == "f1-score":
        score = metrics.f1_score(y_true, y_pred)
    elif metric == "auROC":
        score = metrics.roc_auc_score(y_true, y_pred)
    elif metric == "precision":
        score = metrics.precision_score(y_true, y_pred)
    elif metric == "sensitivity": #same as recall
        score = metrics.recall_score(y_true, y_pred)
    elif metric == "specificity":
        tn, fp, _, _ = metrics.confusion_matrix(y_true, y_label).ravel()
        return tn / (tn + fp)

    return score

    ### ===== TODO : END ===== ###
```

1b)

```
def cv_performance(clf, X, y, kf, metric="accuracy"):
    """
    Splits the data, X and y, into k-folds and runs k-fold cross-validation.
    Trains classifier on k-1 folds and tests on the remaining fold.
    Calculates the k-fold cross-validation performance metric for classifier
    by averaging the performance across folds.

    Parameters
    -----
        clf      -- classifier (instance of LinearSVC)
        X        -- numpy array of shape (n,d), feature vectors
                   n = number of examples
                   d = number of features
        y        -- numpy array of shape (n,), binary labels {1,-1}
        kf       -- model_selection.StratifiedKFold
        metric   -- string, option used to select performance measure

    Returns
    -----
        score    -- float, average cross-validation performance across k folds
    """

    ### ===== TODO : START ===== ###
    # part 1b: compute average cross-validation performance
    scores = []

    for train_index, test_index in kf.split(X, y):
        X_train, X_test = X[train_index], X[test_index]
        y_train, y_test = y[train_index], y[test_index]

        clf.fit(X_train, y_train)

        y_scores = clf.decision_function(X_test)

        if metric == "auROC":
            score = performance(y_test, y_scores, metric)
        else:
            y_pred = np.sign(y_scores)
            y_pred[y_pred == 0] = 1
            score = performance(y_test, y_pred, metric)

        scores.append(score)

    return np.mean(scores)
    ### ===== TODO : END ===== ###
```



(c)

```
def select_param_linear(X, y, kf, metric="accuracy"):
    """
    Sweeps different settings for the hyperparameter of a linear SVM,
    calculating the k-fold CV performance for each setting, then selecting the
    hyperparameter that 'maximize' the average k-fold CV performance.

    Parameters
    -----
    X      -- numpy array of shape (n,d), feature vectors
            |      |      |
            |      |      |      n = number of examples
            |      |      |      d = number of features
    y      -- numpy array of shape (n,), binary labels {1,-1}
    kf     -- model_selection.StratifiedKFold
    metric -- string, option used to select performance measure

    Returns
    -----
    C -- float, optimal parameter value for linear SVM
    """

    print('Linear SVM Hyperparameter Selection based on ' + str(metric) + ':')
    C_range = 10.0 ** np.arange(-3, 3)

    ### ===== TODO : START ===== ###
    # part 1c: select optimal hyperparameter using cross-validation

    best_score = 0
    best_C = None

    for C in C_range:
        clf = LinearSVC(loss = 'hinge', random_state=0, C=C)
        score = cv_performance(clf, X, y, kf, metric)
        if score > best_score:
            best_score = score
            best_C = C

    print("For {}: Best C = {:.3f}. Best Score = {:.3f}".format(metric, best_C, best_score))

    return best_C
    ### ===== TODO : END ===== ###
```

3.2.(b).

```
def performance_test(clf, X, y, metric="accuracy"):
    """
    Estimates the performance of the classifier.

    Parameters
    -----
    |   clf           |   -- classifier (instance of LinearSVC)
    |   |   |   |   |   |   [already fit to data]
    |   X           |   -- numpy array of shape (n,d), feature vectors of test set
    |   |   |   |   |   |   n = number of examples
    |   |   |   |   |   |   d = number of features
    |   y           |   -- numpy array of shape (n,), binary labels {1,-1} of test set
    |   metric      |   -- string, option used to select performance measure

    Returns
    -----
    |   score        |   -- float, classifier performance
    |   |   |   |   |   |
    """

    ### ===== TODO : START ===== ###
    # part 2b: return performance on test data under a metric.
    y_scores = clf.decision_function(X)

    if metric != "auroc":
        |   y_pred = np.sign(y_scores)
        |   y_pred[y_pred == 0] = 1
    else:
        |   y_pred = y_scores

    # Use the performance function to calculate the score
    score = performance(y, y_pred, metric)
    return score

    ### ===== TODO : END ===== ###
```

3.1.(b), (c) , 3.2.(a), (b) . main function:

```
#####  
# main  
#####  
  
def main() :  
    np.random.seed(1234)  
  
    # read the tweets and its labels, change the following two lines to your own path.  
    ### ===== TODO : START ===== ###  
    file_path = '../data/tweets.txt'  
    label_path = '../data/labels.txt'  
    ### ===== TODO : END ===== ###  
    dictionary = extract_dictionary(file_path)  
    print(len(dictionary))  
    X = extract_feature_vectors(file_path, dictionary)  
    y = read_vector_file(label_path)  
    # split data into training (training + cross-validation) and testing set  
    X_train, X_test = X[:560], X[560:]  
    y_train, y_test = y[:560], y[560:]  
  
    metric_list = ["accuracy", "f1-score", "auroc", "precision", "sensitivity", "specificity"]  
  
    ### ===== TODO : START ===== ###  
    # part 1b: create stratified folds (5-fold CV)  
    kf = StratifiedKFold(n_splits = 5, shuffle = True, random_state = 0)  
  
    # part 1c: for each metric, select optimal hyperparameter for linear SVM using CV  
    best_C_values = {}  
    for metric in metric_list:  
        best_C_values[metric] = select_param_linear(X_train, y_train, kf, metric)  
  
    # part 2a: train linear SVMs with selected hyperparameters  
    classifiers = {}  
    for metric, C in best_C_values.items():  
        clf = LinearSVC(loss='hinge', random_state=0, C=C)  
        clf.fit(X_train, y_train)  
        classifiers[metric] = clf  
  
    # part 2b: test the performance of your classifiers.  
    for metric, clf in classifiers.items():  
        score = performance_test(clf, X_test, y_test, metric)  
        print(f"Performance for {metric} with C={C}: {score}")  
    ### ===== TODO : END ===== ###  
  
if __name__ == "__main__" :  
    main()
```

```
Linear SVM Hyperparameter Selection based on accuracy:  
For accuracy: Best C = 10.000. Best Score = 0.829  
Linear SVM Hyperparameter Selection based on f1-score:  
For f1-score: Best C = 10.000. Best Score = 0.883  
Linear SVM Hyperparameter Selection based on auroc:  
For auroc: Best C = 10.000. Best Score = 0.895  
Linear SVM Hyperparameter Selection based on precision:  
For precision: Best C = 10.000. Best Score = 0.856  
Linear SVM Hyperparameter Selection based on sensitivity:  
For sensitivity: Best C = 0.001. Best Score = 1.000  
Linear SVM Hyperparameter Selection based on specificity:  
For specificity: Best C = 10.000. Best Score = 0.625  
Performance for accuracy with C=10.0: 0.7428571428571429  
Performance for f1-score with C=10.0: 0.43749999999999994  
Performance for auroc with C=10.0: 0.7453838678328474  
Performance for precision with C=10.0: 0.6363636363636364  
Performance for sensitivity with C=10.0: 1.0  
Performance for specificity with C=10.0: 0.9183673469387755
```



## Problem 4:

(a).

```
### ===== TODO : START ===== ###
# Part 4(a): Implement the decision tree classifier and report the training error.
print('Classifying using Decision Tree...')
clf = DecisionTreeClassifier(criterion='entropy', random_state=0)
clf.fit(X, y)
ypred_train = clf.predict(X)
train_error = 1 - metrics.accuracy_score(y, ypred_train)
print(f"Training Error: {train_error}")
train_error, test_error = error(clf, X, y)
print(f"Average Training Error after 100 trials: {train_error}")
print(f"Average Test Error after 100 trials: {test_error}")
### ===== TODO : END ===== ###
```

✓ 0.3s

```
Classifying using Decision Tree...
Training Error: 0.014044943820224698
Average Training Error after 100 trials: 0.011528998242530775
Average Test Error after 100 trials: 0.24104895104895108
```

(b)

```
### ===== TODO : START ===== ###
# Part 4(b): Implement the random forest classifier and adjust the number of samples used in bootstrap sampling.
print('Classifying using Random Forest...')
n_samples = len(X) # Total number of samples in your data
max_samples_options = [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8]

best_test_error = float('inf')
best_max_samples = None
best_train_error = None

for max_samples_ratio in max_samples_options:
    max_samples = int(n_samples * max_samples_ratio)
    clf = RandomForestClassifier(criterion='entropy', random_state=0, max_samples=max_samples)

    # Use the provided error function to get training and test error
    train_error, test_error = error(clf, X, y)

    if test_error < best_test_error:
        best_test_error = test_error
        best_train_error = train_error
        best_max_samples = max_samples

print(f"Best setting for max_samples: {best_max_samples} samples")
print(f"Training Error for this setting: {best_train_error}")
print(f"Test Error for this setting: {best_test_error}")
### ===== TODO : END ===== ###
```

✓ 1m 52.0s

```
Classifying using Random Forest...
Best setting for max_samples: 142 samples
Training Error for this setting: 0.10314586994727591
Test Error for this setting: 0.18797202797202794
```



(c)

```
### ===== TODO : START ===== ###
# Part 4(c): Implement the random forest classifier and adjust the number of features for each decision tree.
print('Classifying using Random Forest...')
best_max_samples = 142

best_test_error = float('inf')
best_max_features = None
best_train_error = None

for max_features in range(1, 8):
    clf = RandomForestClassifier(criterion='entropy', random_state=0, max_samples=best_max_samples, max_features=max_features)

    train_error, test_error = error(clf, X, y)

    if test_error < best_test_error:
        best_test_error = test_error
        best_train_error = train_error
        best_max_features = max_features

print(f"Best setting for max_features: {best_max_features}")
print(f"Training Error for this setting: {best_train_error}")
print(f"Test Error for this setting: {best_test_error}")
### ===== TODO : END ===== ###
```

✓ 1m 29.9s

```
Classifying using Random Forest...
Best setting for max_features: 3
Training Error for this setting: 0.10244288224956065
Test Error for this setting: 0.1872727272727273
```

```
In [ ]: import os
import sys

In [ ]: # To add your own Drive Run this cell.
from google.colab import drive
drive.mount('/content/drive')

In [ ]: # Please append your own directory after '/content/drive/My Drive/'
### ===== TODO : START ===== ###
sys.path += ['/content/drive/My Drive/path_to_your_code']
### ===== TODO : END ===== ###

In [ ]: """
Author      : Yi-Chieh Wu, Sriram Sankararman
Description : Twitter
"""

from string import punctuation

import numpy as np
import matplotlib.pyplot as plt
# !!! MAKE SURE TO USE LinearSVC.decision_function(X), NOT LinearSVC.predict(X)
# (this makes 'continuous-valued' predictions)
from sklearn.svm import LinearSVC
from sklearn.model_selection import StratifiedKFold
from sklearn import metrics
```

## Problem 3: Twitter Analysis Using SVM

```
In [ ]: #####
# functions -- input/output
#####

def read_vector_file(fname):
    """
    Reads and returns a vector from a file.

    Parameters
    -----
        fname -- string, filename

    Returns
    -----
        labels -- numpy array of shape (n,)
                n is the number of non-blank lines in the text file
    """
    return np.genfromtxt(fname)

def write_label_answer(vec, outfile):
    """
    Writes your label vector to the given file.

    Parameters
    -----
        vec      -- numpy array of shape (n,) or (n,1), predicted scores
        outfile -- string, output filename
    """

    # for this project, you should predict 70 labels
    if(vec.shape[0] != 70):
        print("Error - output vector should have 70 rows.")
        print("Aborting write.")
        return

    np.savetxt(outfile, vec)
```

```
In [ ]: #####
# functions -- feature extraction
#####

def extract_words(input_string):
    """
    Processes the input_string, separating it into "words" based on the p
    of spaces, and separating punctuation marks into their own words.

    Parameters
    -----
        input_string -- string of characters

    Returns
    -----
        words      -- list of lowercase "words"
    """
```

```

for c in punctuation :
    input_string = input_string.replace(c, ' ' + c + ' ')
return input_string.lower().split()

def extract_dictionary(infile):
    """
    Given a filename, reads the text file and builds a dictionary of unique
    words/punctuations.

    Parameters
    -----
        infile      -- string, filename

    Returns
    -----
        word_list -- dictionary, (key, value) pairs are (word, index)
    """

    word_list = {}
    idx = 0
    with open(infile, 'r') as fid :
        # process each line to populate word_list
        for input_string in fid:
            words = extract_words(input_string)
            for word in words:
                if word not in word_list:
                    word_list[word] = idx
                    idx += 1
    return word_list

def extract_feature_vectors(infile, word_list):
    """
    Produces a bag-of-words representation of a text file specified by the
    filename infile based on the dictionary word_list.

    Parameters
    -----
        infile      -- string, filename
        word_list    -- dictionary, (key, value) pairs are (word, index)

    Returns
    -----
        feature_matrix -- numpy array of shape (n,d)
                        boolean (0,1) array indicating word presence in
                        n is the number of non-blank lines in the text
                        d is the number of unique words in the text file
    """

    num_lines = sum(1 for line in open(infile, 'r'))
    num_words = len(word_list)
    feature_matrix = np.zeros((num_lines, num_words))

    with open(infile, 'r') as fid :
        # process each line to populate feature_matrix

```



```

    for i, input_string in enumerate(fid):
        words = extract_words(input_string)
        for word in words:
            feature_matrix[i, word_list[word]] = 1.0

    return feature_matrix

```

```

In [ ]: #####
# functions -- evaluation
#####

def performance(y_true, y_pred, metric="accuracy"):
    """
    Calculates the performance metric based on the agreement between the
    true labels and the predicted labels.

    Parameters
    -----
        y_true -- numpy array of shape (n,), known labels
        y_pred -- numpy array of shape (n,), (continuous-valued) predictions
        metric -- string, option used to select the performance measure
                  options: 'accuracy', 'f1-score', 'auroc', 'precision',
                           'sensitivity', 'specificity'

    Returns
    -----
        score -- float, performance score
    """
    # map continuous-valued predictions to binary labels
    y_label = np.sign(y_pred)
    y_label[y_label==0] = 1

    ### ===== TODO : START ===== ###
    # part 1a: compute classifier performance

    if metric == "accuracy":
        score = metrics.accuracy_score(y_true, y_pred)
    elif metric == "f1-score":
        score = metrics.f1_score(y_true, y_pred)
    elif metric == "auroc":
        score = metrics.roc_auc_score(y_true, y_pred)
    elif metric == "precision":
        score = metrics.precision_score(y_true, y_pred)
    elif metric == "sensitivity": #same as recall
        score = metrics.recall_score(y_true, y_pred)
    elif metric == "specificity":
        tn, fp, _, _ = metrics.confusion_matrix(y_true, y_label).ravel()
        return tn / (tn + fp)

    return score
    ### ===== TODO : END ===== ###

def cv_performance(clf, X, y, kf, metric="accuracy"):
    """
    Splits the data, X and y, into k-folds and runs k-fold cross-validation

```

Trains classifier on k-1 folds and tests on the remaining fold. Calculates the k-fold cross-validation performance metric for classification by averaging the performance across folds.

#### Parameters

-----

```
clf      -- classifier (instance of LinearSVC)
X        -- numpy array of shape (n,d), feature vectors
           n = number of examples
           d = number of features
y        -- numpy array of shape (n,), binary labels {1,-1}
kf       -- model_selection.StratifiedKFold
metric   -- string, option used to select performance measure
```

#### Returns

-----

```
score    -- float, average cross-validation performance across k folds
"""
```

```
### ===== TODO : START ===== ###
```

```
# part 1b: compute average cross-validation performance
```

```
scores = []
```

```
for train_index, test_index in kf.split(X, y):
    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]

    clf.fit(X_train, y_train)

    y_scores = clf.decision_function(X_test)

    if metric == "auroc":
        score = performance(y_test, y_scores, metric)
    else:
        y_pred = np.sign(y_scores)
        y_pred[y_pred == 0] = 1
        score = performance(y_test, y_pred, metric)

    scores.append(score)
```

```
return np.mean(scores)
```

```
### ===== TODO : END ===== ###
```

```
def select_param_linear(X, y, kf, metric="accuracy"):
    """
```

Sweeps different settings for the hyperparameter of a linear SVM, calculating the k-fold CV performance for each setting, then selecting the hyperparameter that 'maximize' the average k-fold CV performance.

#### Parameters

-----

```
X        -- numpy array of shape (n,d), feature vectors
           n = number of examples
           d = number of features
y        -- numpy array of shape (n,), binary labels {1,-1}
kf       -- model_selection.StratifiedKFold
```

```

        metric -- string, option used to select performance measure

Returns
-----
    C -- float, optimal parameter value for linear SVM
    """

print('Linear SVM Hyperparameter Selection based on ' + str(metric) +
      C_range = 10.0 ** np.arange(-3, 3))

### ===== TODO : START ===== ###
# part 1c: select optimal hyperparameter using cross-validation

best_score = 0
best_C = None

for C in C_range:
    clf = LinearSVC(loss = 'hinge', random_state=0, C=C)
    score = cv_performance(clf, X, y, kf, metric)
    if score > best_score:
        best_score = score
        best_C = C

print("For {}: Best C = {:.3f}. Best Score = {:.3f}".format(metric, b

return best_C
### ===== TODO : END ===== ###

def performance_test(clf, X, y, metric="accuracy"):
    """
    Estimates the performance of the classifier.

    Parameters
    -----
        clf          -- classifier (instance of LinearSVC)
                       [already fit to data]
        X            -- numpy array of shape (n,d), feature vectors of te
                       n = number of examples
                       d = number of features
        y            -- numpy array of shape (n,), binary labels {1,-1} o
        metric       -- string, option used to select performance measure

    Returns
    -----
        score        -- float, classifier performance
    """

    ### ===== TODO : START ===== ###
    # part 2b: return performance on test data under a metric.
    y_scores = clf.decision_function(X)

    if metric != "auroc":
        y_pred = np.sign(y_scores)
        y_pred[y_pred == 0] = 1
    else:

```

```

        y_pred = y_scores

    # Use the performance function to calculate the score
    score = performance(y, y_pred, metric)
    return score

### ===== TODO : END ===== ###

```

```

In [ ]: #####
# main
#####

def main() :
    np.random.seed(1234)

    # read the tweets and its labels, change the following two lines to y
    ### ===== TODO : START ===== ###
    file_path = '../data/tweets.txt'
    label_path = '../data/labels.txt'
    ### ===== TODO : END ===== ###
    dictionary = extract_dictionary(file_path)
    print(len(dictionary))
    X = extract_feature_vectors(file_path, dictionary)
    y = read_vector_file(label_path)
    # split data into training (training + cross-validation) and testing
    X_train, X_test = X[:560], X[560:]
    y_train, y_test = y[:560], y[560:]

    metric_list = ["accuracy", "f1-score", "auroc", "precision", "sensitivity"]

    ### ===== TODO : START ===== ###
    # part 1b: create stratified folds (5-fold CV)
    kf = StratifiedKFold(n_splits = 5, shuffle = True, random_state = 0)

    # part 1c: for each metric, select optimal hyperparameter for linear
    best_C_values = {}
    for metric in metric_list:
        best_C_values[metric] = select_param_linear(X_train, y_train, kf, metric)

    # part 2a: train linear SVMs with selected hyperparameters
    classifiers = {}
    for metric, C in best_C_values.items():
        clf = LinearSVC(loss='hinge', random_state=0, C=C)
        clf.fit(X_train, y_train)
        classifiers[metric] = clf

    # part 2b: test the performance of your classifiers.
    for metric, clf in classifiers.items():
        score = performance_test(clf, X_test, y_test, metric)
        print(f"Performance for {metric} with C={C}: {score}")
    ### ===== TODO : END ===== ###

if __name__ == "__main__" :
    main()

```



```

1811
Linear SVM Hyperparameter Selection based on accuracy:
For accuracy: Best C = 10.000. Best Score = 0.829
Linear SVM Hyperparameter Selection based on f1-score:
For f1-score: Best C = 10.000. Best Score = 0.883
Linear SVM Hyperparameter Selection based on auroc:
For auroc: Best C = 10.000. Best Score = 0.895
Linear SVM Hyperparameter Selection based on precision:
For precision: Best C = 10.000. Best Score = 0.856
Linear SVM Hyperparameter Selection based on sensitivity:
For sensitivity: Best C = 0.001. Best Score = 1.000
Linear SVM Hyperparameter Selection based on specificity:
For specificity: Best C = 10.000. Best Score = 0.625
Performance for accuracy with C=10.0: 0.7428571428571429
Performance for f1-score with C=10.0: 0.43749999999999994
Performance for auroc with C=10.0: 0.7453838678328474
Performance for precision with C=10.0: 0.6363636363636364
Performance for sensitivity with C=10.0: 1.0
Performance for specificity with C=10.0: 0.9183673469387755

```

## Problem 4: Boosting vs. Decision Tree

```

In [ ]: from sklearn.tree import DecisionTreeClassifier
        from sklearn.ensemble import RandomForestClassifier
        from sklearn import metrics
        from sklearn.model_selection import cross_val_score, train_test_split

```

```

In [ ]: class Data :

        def __init__(self) :
            """
            Data class.

            Attributes
            -----
                X -- numpy array of shape (n,d), features
                y -- numpy array of shape (n,), targets
            """

            # n = number of examples, d = dimensionality
            self.X = None
            self.y = None

            self.Xnames = None
            self.yname = None

        def load(self, filename, header=0, predict_col=-1) :
            """Load csv file into X array of features and y array of labels."""

            # determine filename
            f = filename

            # load data
            with open(f, 'r') as fid :

```

```

        data = np.loadtxt(fid, delimiter=",", skiprows=header)

    # separate features and labels
    if predict_col is None :
        self.X = data[:, :]
        self.y = None
    else :
        if data.ndim > 1 :
            self.X = np.delete(data, predict_col, axis=1)
            self.y = data[:, predict_col]
        else :
            self.X = None
            self.y = data[:, :]

    # load feature and label names
    if header != 0:
        with open(f, 'r') as fid :
            header = fid.readline().rstrip().split(",")

        if predict_col is None :
            self.Xnames = header[: ]
            self.ynames = None
        else :
            if len(header) > 1 :
                self.Xnames = np.delete(header, predict_col)
                self.ynames = header[predict_col]
            else :
                self.Xnames = None
                self.ynames = header[0]
    else:
        self.Xnames = None
        self.ynames = None

    # helper functions
    def load_data(filename, header=0, predict_col=-1) :
        """Load csv file into Data class."""
        data = Data()
        data.load(filename, header=header, predict_col=predict_col)
        return data

```

```

In [ ]: # Change the path to your own data directory
        ### ===== TODO : START ===== ###
        titanic = load_data("../data/titanic_train.csv", header=1, predict_col=0)
        ### ===== TODO : END ===== ###
        X = titanic.X; Xnames = titanic.Xnames
        y = titanic.y; yname = titanic.ynames
        n,d = X.shape # n = number of examples, d = number of features

```

```
In [ ]: def error(clf, X, y, ntrials=100, test_size=0.2) :
        """
        Computes the classifier error over a random split of the data,
        averaged over ntrials runs.

        Parameters
        -----
        clf          -- classifier
        X            -- numpy array of shape (n,d), features values
        y            -- numpy array of shape (n,), target classes
        ntrials      -- integer, number of trials
        test_size    -- proportion of data used for evaluation

        Returns
        -----
        train_error  -- float, training error
        test_error   -- float, test error
        """

        train_error = 0
        test_error = 0

        train_scores = []; test_scores = [];
        for i in range(ntrials):
            xtrain, xtest, ytrain, ytest = train_test_split (X,y, test_size =
            clf.fit (xtrain, ytrain)

            ypred = clf.predict (xtrain)
            err = 1 - metrics.accuracy_score (ytrain, ypred, normalize = True)
            train_scores.append (err)

            ypred = clf.predict (xtest)
            err = 1 - metrics.accuracy_score (ytest, ypred, normalize = True)
            test_scores.append (err)

        train_error = np.mean (train_scores)
        test_error = np.mean (test_scores)
        return train_error, test_error
```

```
In [ ]: ### ===== TODO : START ===== ###
        # Part 4(a): Implement the decision tree classifier and report the traini
        print('Classifying using Decision Tree...')
        clf = DecisionTreeClassifier(criterion='entropy', random_state=0)
        clf.fit(X, y)
        ypred_train = clf.predict(X)
        train_error = 1 - metrics.accuracy_score(y, ypred_train)
        print(f"Training Error: {train_error}")
        train_error, test_error = error(clf, X, y)
        print(f"Average Training Error after 100 trials: {train_error}")
        print(f"Average Test Error after 100 trials: {test_error}")
        ### ===== TODO : END ===== ###
```

```
Classifying using Decision Tree...
Training Error: 0.014044943820224698
Average Training Error after 100 trials: 0.011528998242530775
Average Test Error after 100 trials: 0.24104895104895108
```

```
In [ ]: ### ===== TODO : START ===== ###
# Part 4(b): Implement the random forest classifier and adjust the number
print('Classifying using Random Forest...')
n_samples = len(X) # Total number of samples in your data
max_samples_options = [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8]

best_test_error = float('inf')
best_max_samples = None
best_train_error = None

for max_samples_ratio in max_samples_options:
    max_samples = int(n_samples * max_samples_ratio)
    clf = RandomForestClassifier(criterion='entropy', random_state=0, max

    # Use the provided error function to get training and test error
    train_error, test_error = error(clf, X, y)

    if test_error < best_test_error:
        best_test_error = test_error
        best_train_error = train_error
        best_max_samples = max_samples

print(f"Best setting for max_samples: {best_max_samples} samples")
print(f"Training Error for this setting: {best_train_error}")
print(f"Test Error for this setting: {best_test_error}")
### ===== TODO : END ===== ###
```

```
Classifying using Random Forest...
Best setting for max_samples: 142 samples
Training Error for this setting: 0.10314586994727591
Test Error for this setting: 0.18797202797202794
```

```
In [ ]: ### ===== TODO : START ===== ###
# Part 4(c): Implement the random forest classifier and adjust the number
print('Classifying using Random Forest...')
best_max_samples = 142

best_test_error = float('inf')
best_max_features = None
best_train_error = None

for max_features in range(1, 8):
    clf = RandomForestClassifier(criterion='entropy', random_state=0, max

    train_error, test_error = error(clf, X, y)

    if test_error < best_test_error:
        best_test_error = test_error
        best_train_error = train_error
        best_max_features = max_features

print(f"Best setting for max_features: {best_max_features}")
print(f"Training Error for this setting: {best_train_error}")
print(f"Test Error for this setting: {best_test_error}")
### ===== TODO : END ===== ###
```



```
Classifying using Random Forest...  
Best setting for max_features: 3  
Training Error for this setting: 0.10244288224956065  
Test Error for this setting: 0.1872727272727273
```